

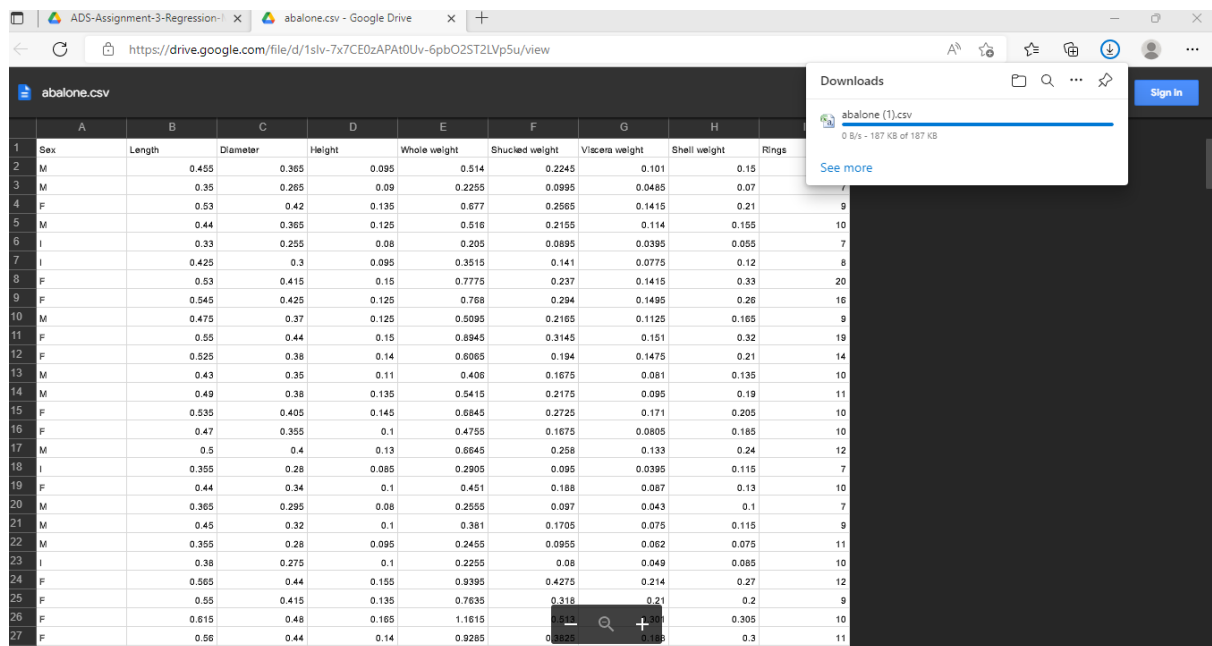
Assignment -3

Regression Model

Assignment Date	29 September 2022
Student Name	Logeshkumar R
Student Roll Number	727719EUCS074
Maximum Marks	2 Marks

Question-1:

Download the dataset: Dataset



	A	B	C	D	E	F	G	H	
1	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
2	M	0.455	0.365	0.095	0.514	0.2245	0.101	0.15	
3	M	0.35	0.265	0.09	0.2255	0.0995	0.0485	0.07	
4	F	0.53	0.42	0.135	0.677	0.2565	0.1415	0.21	
5	M	0.44	0.365	0.125	0.516	0.2155	0.114	0.155	10
6	I	0.33	0.255	0.08	0.205	0.0895	0.0395	0.055	7
7	I	0.425	0.3	0.095	0.3515	0.141	0.0775	0.12	8
8	F	0.53	0.415	0.15	0.7775	0.237	0.1415	0.33	20
9	F	0.545	0.425	0.125	0.768	0.294	0.1495	0.26	16
10	M	0.475	0.37	0.125	0.5095	0.2165	0.1125	0.165	9
11	F	0.55	0.44	0.15	0.8945	0.3145	0.151	0.32	19
12	F	0.525	0.38	0.14	0.6065	0.194	0.1475	0.21	14
13	M	0.43	0.35	0.11	0.406	0.1675	0.081	0.135	10
14	M	0.49	0.38	0.135	0.5415	0.2175	0.095	0.19	11
15	F	0.535	0.405	0.145	0.6645	0.2725	0.171	0.205	10
16	F	0.47	0.355	0.1	0.4755	0.1675	0.0805	0.165	10
17	M	0.5	0.4	0.13	0.6645	0.258	0.133	0.24	12
18	I	0.355	0.28	0.085	0.2905	0.095	0.0395	0.115	7
19	F	0.44	0.34	0.1	0.451	0.188	0.087	0.13	10
20	M	0.365	0.295	0.08	0.2555	0.097	0.043	0.1	7
21	M	0.45	0.32	0.1	0.381	0.1705	0.075	0.115	9
22	M	0.355	0.28	0.095	0.2455	0.0955	0.062	0.075	11
23	I	0.38	0.275	0.1	0.2255	0.08	0.049	0.085	10
24	F	0.565	0.44	0.155	0.9395	0.4275	0.214	0.27	12
25	F	0.55	0.415	0.135	0.7635	0.318	0.21	0.2	9
26	F	0.615	0.48	0.165	1.1615	0.513	0.299	0.305	10
27	F	0.56	0.44	0.14	0.9285	0.4625	0.249	0.3	11

Question-2:

Load the dataset.

Solution:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

d = pd.read_csv("E://abalone (1).csv")

d.head()
```

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
d = pd.read_csv("E://abalone (1).csv")
d.head()
```

```
Out[2]:
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7

Question-3:


Perform Below Visualizations.

· Univariate Analysis

Solution:

```
In [5]: sns.barplot(d.Diameter)
```

```
Out[5]: <AxesSubplot:xlabel='Diameter'>
```




- Bi-Variate Analysis

Solution:

```
In [9]: sns.boxplot(y=d.Rings,x=d.Height)|
```

```
Out[9]: <AxesSubplot:xlabel='Height', ylabel='Rings'>
```

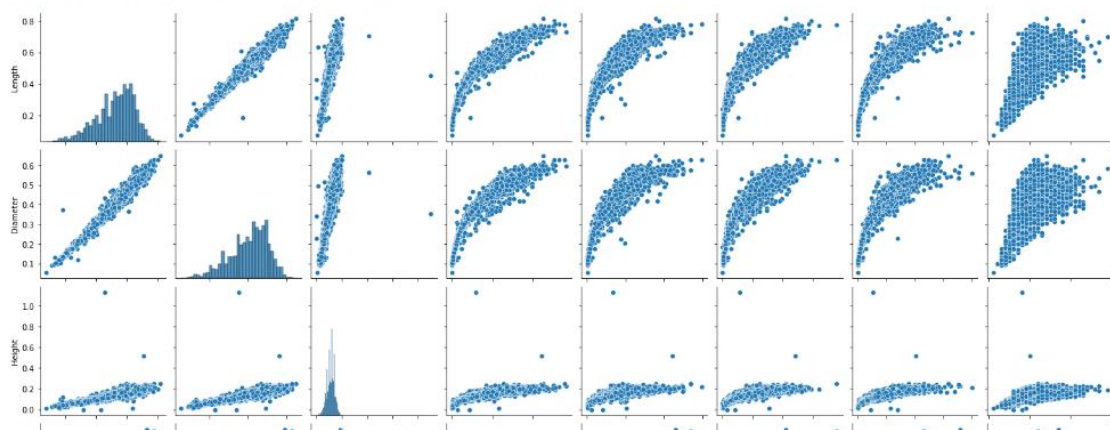


- Multi-Variate Analysis

Solution:

```
In [8]: sns.pairplot(d)
```

```
Out[8]: <seaborn.axisgrid.PairGrid at 0x2cb4763d400>
```



Question-4:

Perform descriptive statistics on the dataset.

Solution:

```
In [10]: d['Diameter'].mean()
```

```
Out[10]: 0.407881254488869
```

```
In [11]: d['Height'].median()
```

```
Out[11]: 0.14
```

```
In [13]: d['Rings'].mode()
```

```
Out[13]: 0    9  
dtype: int64
```

```
In [14]: d.skew()
```

```
Out[14]: Length      -0.639873  
Diameter    -0.609198  
Height       3.128817  
Whole weight  0.530959  
Shucked weight  0.719098  
Viscera weight  0.591852  
Shell weight  0.620927  
Rings        1.114102  
dtype: float64
```

```
In [15]: d.kurt()
```

```
Out[15]: Length      0.064621  
Diameter    -0.045476  
Height      76.025509  
Whole weight -0.023644  
Shucked weight  0.595124  
Viscera weight  0.084012  
Shell weight  0.531926  
Rings        2.330687  
dtype: float64
```

Question-5:

Check for Missing values and deal with them.

Solution:

```
In [16]: d.isna().any()
```

```
Out[16]: Sex           False
Length        False
Diameter       False
Height         False
Whole weight   False
Shucked weight False
Viscera weight False
Shell weight   False
Rings          False
dtype: bool
```

```
In [18]: d['Rings'].fillna(d['Rings'].mean(),inplace=True)
d
```

```
Out[18]:
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.1500	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.0700	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.2100	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.1550	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.0550	7
...
4172	F	0.565	0.450	0.165	0.8870	0.3700	0.2390	0.2490	11


Question-6:

Find the outliers and replace the outliers

Solution:

```
In [19]: sns.boxplot(d['Diameter'])
```

```
Out[19]: <AxesSubplot:xlabel='Diameter'>
```



```
In [20]: Q1=d.Diameter.quantile(0.25)
Q2=d.Diameter.quantile(0.75)
IQR=Q2-Q1
print(IQR)
```

```
0.13
```

```
In [21]: d[d[~((d.Diameter<(Q1-1.5*IQR))|(d.Diameter>(Q2+1.5*IQR)))]]
d
```

Out[21]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.1500	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.0700	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.2100	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.1550	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.0550	7
...
4172	F	0.565	0.450	0.165	0.8870	0.3700	0.2390	0.2490	11
4173	M	0.590	0.440	0.135	0.9660	0.4390	0.2145	0.2605	10
4174	M	0.600	0.475	0.205	1.1760	0.5255	0.2875	0.3080	9
4175	F	0.625	0.485	0.150	1.0945	0.5310	0.2610	0.2960	10
4176	M	0.710	0.555	0.195	1.9485	0.9455	0.3765	0.4950	12

4118 rows × 9 columns

Question-7:

Check for Categorical columns and perform encoding.

Solution:

```
In [22]: d['Sex'].replace({'M':1, 'F':0},inplace=True)
d.head()
```

Out[22]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	1	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	1	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	0	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	1	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	1	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7

Question-8:

Split the data into dependent and independent variables.

Solution:

```
In [23]: dm = pd.get_dummies(d, columns=['Height'])
dm
```

Out[23]:

	Sex	Length	Diameter	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings	Height_0.0	Height_0.015	...	Height_0.21	Height_0.215	Height_0.22	Height_0.225	Height_0.23
0	1	0.455	0.365	0.5140	0.2245	0.1010	0.1500	15	0	0	...	0	0	0	0	0
1	1	0.350	0.265	0.2255	0.0995	0.0485	0.0700	7	0	0	...	0	0	0	0	0
2	0	0.530	0.420	0.6770	0.2565	0.1415	0.2100	9	0	0	...	0	0	0	0	0
3	1	0.440	0.365	0.5160	0.2155	0.1140	0.1550	10	0	0	...	0	0	0	0	0
4	1	0.330	0.255	0.2050	0.0895	0.0395	0.0550	7	0	0	...	0	0	0	0	0
...
4172	0	0.565	0.450	0.8870	0.3700	0.2390	0.2490	11	0	0	...	0	0	0	0	0
4173	1	0.590	0.440	0.9660	0.4390	0.2145	0.2605	10	0	0	...	0	0	0	0	0
4174	1	0.600	0.475	1.1760	0.5255	0.2875	0.3080	9	0	0	...	0	0	0	0	0
4175	0	0.625	0.485	1.0945	0.5310	0.2610	0.2960	10	0	0	...	0	0	0	0	0
4176	1	0.710	0.555	1.9485	0.9455	0.3765	0.4950	12	0	0	...	0	0	0	0	0

4118 rows x 54 columns

```
In [24]: y = dm['Diameter']
y
```

```
Out[24]: 0      0.365
1      0.265
2      0.420
3      0.365
4      0.255
...
4172    0.450
4173    0.440
4174    0.475
4175    0.485
4176    0.555
Name: Diameter, Length: 4118, dtype: float64
```

```
In [25]: x = dm.drop(columns='Diameter', axis=1)
x.head()
```

Out[25]:

	Sex	Length	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings	Height_0.0	Height_0.015	Height_0.04	...	Height_0.21	Height_0.215	Height_0.22	Height_0.225	Height_0.23
0	1	0.455	0.5140	0.2245	0.1010	0.150	15	0	0	0	...	0	0	0	0	0
1	1	0.350	0.2255	0.0995	0.0485	0.070	7	0	0	0	...	0	0	0	0	0
2	0	0.530	0.6770	0.2565	0.1415	0.210	9	0	0	0	...	0	0	0	0	0
3	1	0.440	0.5160	0.2155	0.1140	0.155	10	0	0	0	...	0	0	0	0	0
4	1	0.330	0.2050	0.0895	0.0395	0.055	7	0	0	0	...	0	0	0	0	0

5 rows x 53 columns

Scale the independent variables

```
In [26]: x=d.iloc[:,6:7].values
          from sklearn.preprocessing import StandardScaler
          std=StandardScaler()
          x=std.fit_transform(x)
          x
```

Question-10:

Split the data into training and testing

```
In [29]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=0)
```

```
Out[30]: array([[ 0.49751414],
                [-0.68803532],
                [-1.2600514 ],
                ...,
                [ 0.55748357],
                [-0.68803532],
                [-0.4481576 ]])
```

```
[ 1.87681099e+00],
[ 8.29652515e-01],
[-1.37076419e+00],
[-1.23237320e+00],
[-2.40571123e-01],
[-6.05000724e-01],
[-1.03862582e+00],
[-2.12892926e-01],
[ 2.02280037e-01],
[ 1.24482548e+00],
[-3.69736045e-01],
[-4.02027276e-01],
[-4.25092440e-01],
```



```
In [32]: y_train
```

```
Out[32]: 780      0.410
          2411     0.395
          1553     0.290
          2627     0.205
          962      0.390
          ...
          1060     0.195
          3312     0.410
          1681     0.540
          2650     0.395
          2777     0.420
          Name: Diameter, Length: 3294, dtype: float64
```

```
In [33]: y_test
```

```
Out[33]: 2004     0.275
          615      0.345
          2888     0.400
          2599     0.480
          464      0.195
          ...
          2102     0.310
          410      0.500
          3138     0.400
          2520     0.425
          3358     0.215
          Name: Diameter, Length: 824, dtype: float64
```

Question-11:

Build the Model

Solution:

```
In [34]: from sklearn.linear_model import LinearRegression
          regressor=LinearRegression()
          regressor.fit(x_train,y_train)
```

```
Out[34]: LinearRegression()
```


Question-12:

Train the Model

Solution:

```
In [35]: plt.scatter(x_train,y_train,color='yellow')  
plt.plot(x_train,regressor.predict(x_train))
```

```
Out[35]: [<matplotlib.lines.Line2D at 0x2cb4c5fadf0>]
```



Question-13:

Test the Model

Solution:

```
In [36]: y_pred=regressor.predict(x_test)  
y_pred
```

```
Out[36]: array([0.30739667, 0.3231634 , 0.37046358, 0.48398401, 0.28374658,  
0.5336492 , 0.44101968, 0.51867081, 0.42643546, 0.56557683,  
0.41027456, 0.43313632, 0.45087389, 0.4272238 , 0.3491785 ,  
0.41421625, 0.43274215, 0.46900562, 0.39884369, 0.38544197,  
0.42682963, 0.5671535 , 0.51985332, 0.35509102, 0.36415689,  
0.4457497 , 0.44377886, 0.45599807, 0.64835214, 0.40475621,  
0.37795277, 0.43116548, 0.43392466, 0.39332533, 0.481619 ,  
0.4619106 , 0.4445672 , 0.32355757, 0.29202411, 0.49738573,  
0.51157579, 0.35469685, 0.38150029, 0.33774762, 0.37125191,  
0.30936751, 0.34563098, 0.36967524, 0.34326597, 0.48634902,  
0.60578198, 0.45718058, 0.5584818 , 0.60420531, 0.44417303,  
0.28216991, 0.37204025, 0.48950237, 0.4556039 , 0.3144917 ,  
0.52300666, 0.29596579, 0.39214283, 0.28492908, 0.42958881,  
0.39293116, 0.47531231, 0.46072809, 0.48437818, 0.35982104,  
0.3677044 , 0.43037714, 0.51551747, 0.53483171, 0.39411367,  
0.28650576, 0.53877339, 0.54823343, 0.45639224, 0.4457497 ,  
0.41500458, 0.39569034, 0.40317954, 0.63495043, 0.5076341 ,
```

Question-14:

Measure the performance using Metrics.

Solution:

```
In [37]: from sklearn.metrics import r2_score  
accuracy=r2_score(y_test,y_pred)  
accuracy
```

```
Out[37]: 0.8137980390529289
```

```
In [38]: from sklearn import metrics  
np.sqrt(metrics.mean_squared_error(y_test,y_pred))
```

```
Out[38]: 0.040880060610600996
```