

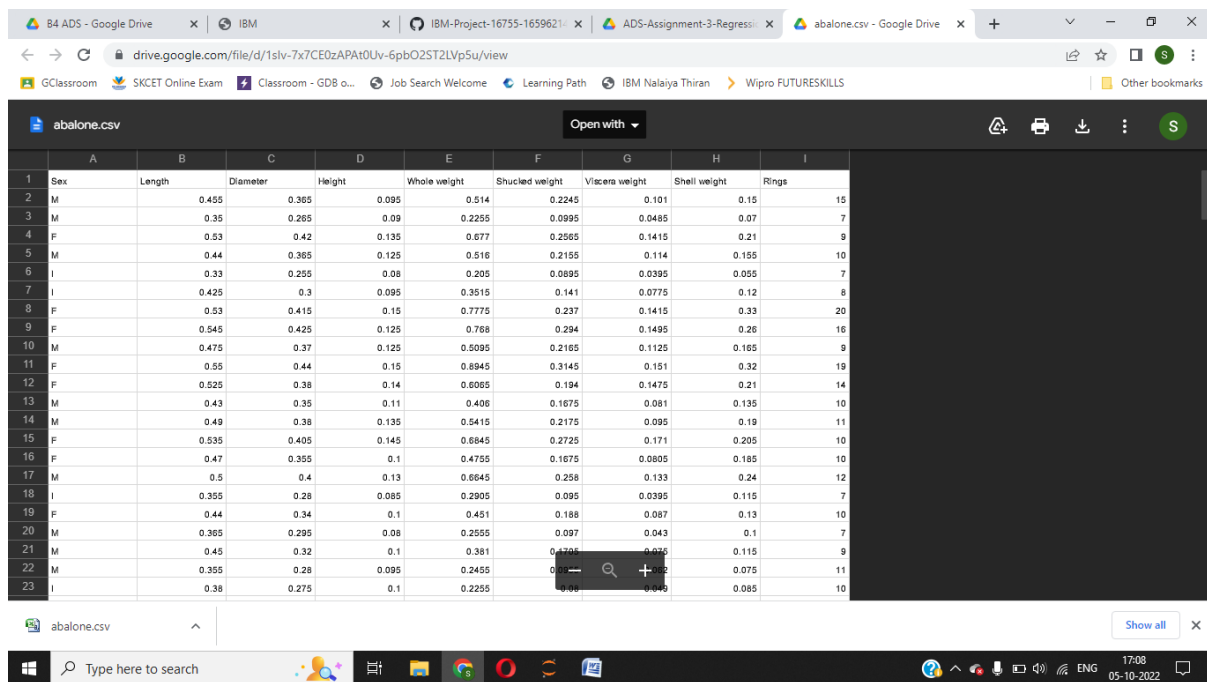
ASSIGNMENT -3

REGRESSION

Assignment Date	29 September 2022
Student Name	Manjunathan V
Student Roll Number	727719EUCS080
Maximum Marks	2 Marks

Question-1:

Download the dataset: Dataset



The screenshot shows a Google Drive interface with a link to the 'abalone.csv' dataset. The dataset is previewed in a spreadsheet format with the following columns: Sex, Length, Diameter, Height, Whole weight, Shucked weight, Viscera weight, Shell weight, and Rings. The data is organized into rows, with the first row being the header and subsequent rows containing numerical values for each attribute.

	A	B	C	D	E	F	G	H	I
1	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
2	M	0.455	0.365	0.095	0.514	0.2245	0.101	0.15	15
3	M	0.35	0.265	0.09	0.2255	0.0995	0.0485	0.07	7
4	F	0.53	0.42	0.135	0.677	0.2565	0.1415	0.21	9
5	M	0.44	0.365	0.125	0.516	0.2155	0.114	0.155	10
6	I	0.33	0.255	0.08	0.205	0.0895	0.0395	0.055	7
7	I	0.425	0.3	0.095	0.3515	0.141	0.0775	0.12	8
8	F	0.53	0.415	0.15	0.7775	0.237	0.1415	0.33	20
9	F	0.545	0.425	0.125	0.768	0.294	0.1495	0.26	16
10	M	0.475	0.37	0.125	0.5095	0.2165	0.1125	0.165	9
11	F	0.55	0.44	0.15	0.8945	0.3145	0.151	0.32	19
12	F	0.525	0.38	0.14	0.6065	0.194	0.1475	0.21	14
13	M	0.43	0.35	0.11	0.406	0.1675	0.081	0.135	10
14	M	0.49	0.38	0.135	0.5415	0.2175	0.095	0.19	11
15	F	0.535	0.405	0.145	0.6845	0.2725	0.171	0.295	10
16	F	0.47	0.355	0.1	0.4755	0.1675	0.0805	0.185	10
17	M	0.5	0.4	0.13	0.6645	0.258	0.133	0.24	12
18	I	0.355	0.28	0.085	0.2905	0.095	0.0395	0.115	7
19	F	0.44	0.34	0.1	0.451	0.188	0.087	0.13	10
20	M	0.365	0.295	0.08	0.2555	0.097	0.043	0.1	7
21	M	0.45	0.32	0.1	0.381	0.1705	0.075	0.115	9
22	M	0.355	0.28	0.095	0.2455	0.102	0.075	0.075	11
23	I	0.38	0.275	0.1	0.2255	0.095	0.045	0.085	10

Question-2:

Load the dataset.

Solution:

```
import pandas as pd

import numpy as np

import matplotlib.pyplot as plt

import seaborn as sns

df = pd.read_csv("C://Users//Mohana Sowdes//Downloads//abalone.csv")

df.head()
```

```
In [3]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
df = pd.read_csv("C://Users//Mohana Sowdesh//Downloads//abalone.csv")
df.head()
```

Out[3]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7

Question-3:


Perform Below Visualizations.

- Univariate Analysis

Solution:

```
In [11]: sns.boxplot(df.Rings)
```

Out[11]: <AxesSubplot:xlabel='Rings'>

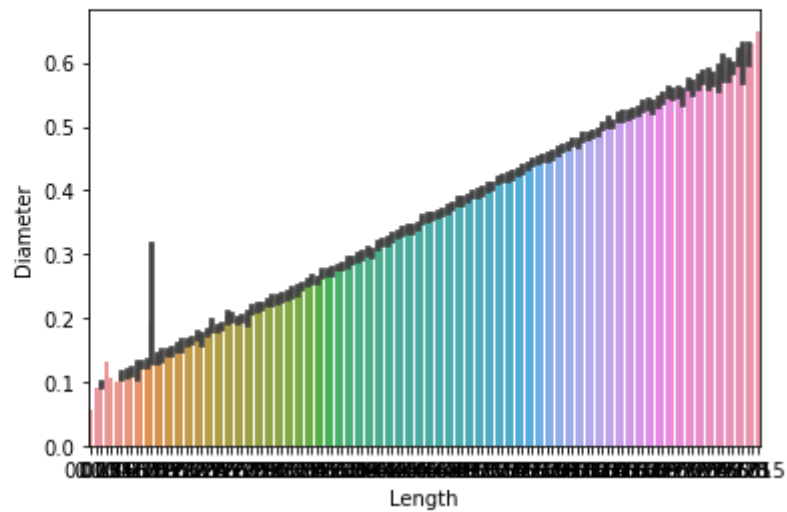


· Bi-Variate Analysis

Solution:

```
In [15]: sns.barplot(y=df.Diameter,x=df.Length)
```

```
Out[15]: <AxesSubplot:xlabel='Length', ylabel='Diameter'>
```




· Multi-Variate Analysis

Solution:

```
In [16]: sns.pairplot(df)
```

```
Out[16]: <seaborn.axisgrid.PairGrid at 0x212ca7ae220>
```



Question-4:

Perform descriptive statistics on the dataset.

Solution:

```
In [17]: df['Rings'].mean()
```

```
Out[17]: 9.933684462532918
```

```
In [21]: df['Length'].median()
```

```
Out[21]: 0.545
```

```
In [22]: df['Sex'].mode()
```

```
Out[22]: 0      M  
dtype: object
```

```
In [23]: df.skew()
```

```
Out[23]: Length          -0.639873  
Diameter          -0.609198  
Height           3.128817  
Whole weight      0.530959  
Shucked weight    0.719098  
Viscera weight    0.591852  
Shell weight      0.620927  
Rings            1.114102  
dtype: float64
```

```
In [24]: df.kurt()
```

```
Out[24]: Length          0.064621  
Diameter          -0.045476  
Height          76.025509  
Whole weight     -0.023644  
Shucked weight    0.595124  
Viscera weight    0.084012  
Shell weight      0.531926  
Rings            2.330687  
dtype: float64
```

```
In [26]: df.var()
```

```
Out[26]: Length          0.014422
Diameter          0.009849
Height          0.001750
Whole weight      0.240481
Shucked weight    0.049268
Viscera weight    0.012015
Shell weight      0.019377
Rings            10.395266
dtype: float64
```

```
In [27]: df.max()
```

```
Out[27]: Sex          M
Length          0.815
Diameter          0.65
Height          1.13
Whole weight     2.8255
Shucked weight    1.488
Viscera weight    0.76
Shell weight      1.005
Rings            29
dtype: object
```

Question-5:

Handle the Missing values.

Solution:

```
In [29]: df.isna().any()
```

```
Out[29]: Sex          False
Length          False
Diameter          False
Height          False
Whole weight     False
Shucked weight    False
Viscera weight    False
Shell weight      False
Rings            False
dtype: bool
```

```
In [30]: df['Height'].fillna(df['Height'].mean(),inplace=True)
df
```

```
Out[30]:
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.1500	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.0700	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.2100	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.1550	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.0550	7
...
4172	F	0.565	0.450	0.165	0.8870	0.3700	0.2390	0.2490	11
4173	M	0.590	0.440	0.135	0.9660	0.4390	0.2145	0.2605	10
4174	M	0.600	0.475	0.205	1.1760	0.5255	0.2875	0.3080	9


Question-6:

Find the outliers and replace the outliers

Solution:

```
In [38]: sns.boxplot(df['Rings'])
```

```
Out[38]: <AxesSubplot:xlabel='Rings'>
```



```
In [40]: Q1=df.Rings.quantile(0.25)
Q2=df.Rings.quantile(0.75)
IQR=Q2-Q1
print(IQR)
```

```
3.0
```

```
In [42]: df=df[~((df.Rings<(Q1-1.5*IQR))|(df.Rings>(Q2+1.5*IQR)))]
df
```

Out[42]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.1500	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.0700	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.2100	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.1550	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.0550	7
...
4172	F	0.565	0.450	0.165	0.8870	0.3700	0.2390	0.2490	11
4173	M	0.590	0.440	0.135	0.9660	0.4390	0.2145	0.2605	10
4174	M	0.600	0.475	0.205	1.1760	0.5255	0.2875	0.3080	9
4175	F	0.625	0.485	0.150	1.0945	0.5310	0.2610	0.2960	10
4176	M	0.710	0.555	0.195	1.9485	0.9455	0.3765	0.4950	12

3899 rows × 9 columns

Question-7:

Check for Categorical columns and perform encoding.

Solution:

```
In [44]: df['Sex'].replace({'F':1,'M':0},inplace=True)
df.head()
```

Out[44]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	0	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	0	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	1	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	0	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	1	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7

Question-8:

Split the data into dependent and independent variables.

Solution:

```
In [45]: data_main= pd.get_dummies(df,columns=['Rings'])
data_main
```

Out[45]:

eight	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings_4	Rings_5	Rings_6	Rings_7	Rings_8	Rings_9	Rings_10	Rings_11	Rings_12	Rings_13	Rings_14	Rings_15
0.095	0.5140	0.2245	0.1010	0.1500	0	0	0	0	0	0	0	0	0	0	0	1
0.090	0.2255	0.0995	0.0485	0.0700	0	0	0	1	0	0	0	0	0	0	0	0
0.135	0.6770	0.2565	0.1415	0.2100	0	0	0	0	0	1	0	0	0	0	0	0
0.125	0.5160	0.2155	0.1140	0.1550	0	0	0	0	0	0	1	0	0	0	0	0
0.080	0.2050	0.0895	0.0395	0.0550	0	0	0	1	0	0	0	0	0	0	0	0
...
0.165	0.8870	0.3700	0.2390	0.2490	0	0	0	0	0	0	0	1	0	0	0	0
0.135	0.9660	0.4390	0.2145	0.2605	0	0	0	0	0	0	1	0	0	0	0	0
0.205	1.1760	0.5255	0.2875	0.3080	0	0	0	0	0	1	0	0	0	0	0	0
0.150	1.0945	0.5310	0.2610	0.2960	0	0	0	0	0	0	1	0	0	0	0	0
0.195	1.9485	0.9455	0.3765	0.4950	0	0	0	0	0	0	0	0	1	0	0	0

```
In [46]: y = data_main['Height']
y
```

Out[46]:

0	0.095
1	0.090
2	0.135
3	0.125
4	0.080
...	...
4172	0.165
4173	0.135
4174	0.205
4175	0.150
4176	0.195

Name: Height, Length: 3899, dtype: float64

```
In [47]: x = data_main.drop(columns='Height',axis=1)
x.head()
```

Out[47]:

	Sex	Length	Diameter	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings_4	Rings_5	Rings_6	Rings_7	Rings_8	Rings_9	Rings_10	Rings_11	Rings_12	Rings_13
0	0	0.455	0.365	0.5140	0.2245	0.1010	0.150	0	0	0	0	0	0	0	0	0	0
1	0	0.350	0.265	0.2255	0.0995	0.0485	0.070	0	0	0	1	0	0	0	0	0	0
2	1	0.530	0.420	0.6770	0.2565	0.1415	0.210	0	0	0	0	0	1	0	0	0	0
3	0	0.440	0.365	0.5160	0.2155	0.1140	0.155	0	0	0	0	0	0	1	0	0	0
4	1	0.330	0.255	0.2050	0.0895	0.0395	0.055	0	0	0	1	0	0	0	0	0	0

Question-9:

Scale the independent variables

Solution:

```
In [48]: x=df.iloc[:,6:7].values
         from sklearn.preprocessing import StandardScaler
         std=StandardScaler()
         x=std.fit_transform(x)
         x
```

```
Out[48]: array([[ -0.69758868],
                [ -1.17989471],
                [ -0.32552403],
                ...,
                [  1.01574608],
                [  0.77229637],
                [  1.83336964]])
```

Question-10:

Split the data into training and testing

Solution:

```
In [50]: from sklearn.model_selection import train_test_split
         x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=0)
```

```
In [51]: x_train
```

```
Out[51]: array([[ -0.75730276],
                [  1.41537108],
                [  0.34511103],
                ...,
                [ -0.04992058],
                [ -1.07884012],
                [  1.02493287]])
```

```
In [52]: y_train
```

```
Out[52]: 3115    0.120
         3626    0.190
         2425    0.170
         822     0.095
         813     0.060
         ...
         974     0.130
         3524    0.085
         1794    0.130
         2820    0.090
         2945    0.190
         Name: Height, Length: 3119, dtype: float64
```



```
In [53]: x_test
```

```
Out[53]: array([[ 0.17974896],  
                [-0.70218207],  
                [-1.40956425],  
                [ 0.3588912 ],  
                [ 0.93765844],  
                [-1.21204845],  
                [-0.99156569],  
                [-0.26580995],  
                [ 0.7585162 ],  
                [ 2.22840125],  
                [ 0.0097935 ],  
                [ 0.27621017],  
                [-0.61490765],  
                [ 1.79662252],  
                [-1.55655276],  
                [-1.30850966],  
                [-1.31310305],  
                [-0.53222661],  
                [ 0.32673746],  
                [ 0.26242899]])
```

```
In [54]: x_test
```

```
Out[54]: array([[ 0.17974896],
                [-0.70218207],
                [-1.40956425],
                [ 0.3588912 ],
                [ 0.93765844],
                [-1.21204845],
                [-0.99156569],
                [-0.26580995],
```

Question-11:

Build the Model

Solution:

```
In [56]: from sklearn.linear_model import LinearRegression
regressor=LinearRegression()
regressor.fit(x_train,y_train)
```

```
Out[56]: LinearRegression()
```


Question-12:

Train the Model

Solution:

```
In [60]: plt.scatter(x_train,y_train,color='violet')
plt.plot(x_train,regressor.predict(x_train))

Out[60]: [<matplotlib.lines.Line2D at 0x212d17166a0>]
```



Question-13:

Test the Model

Solution:

```
In [62]: y_pred=regressor.predict(x_test)
y_pred
```

```
Out[62]: array([0.14371637, 0.11505347, 0.09206344, 0.14953852, 0.16834854,
0.09848273, 0.10564846, 0.12923563, 0.16252639, 0.21029789,
0.13819279, 0.14685137, 0.1178899 , 0.19626501, 0.08728629,
0.09534773, 0.09519844, 0.12057705, 0.14849352, 0.14640351,
0.15491281, 0.10460345, 0.10400631, 0.09415344, 0.12296562,
0.08639057, 0.12729491, 0.10206559, 0.13147492, 0.09221272,
0.1773057 , 0.17118498, 0.1561071 , 0.16715426, 0.15759996,
0.13878993, 0.14341779, 0.1243092 , 0.19193572, 0.14341779,
0.09728844, 0.12177134, 0.19656358, 0.16819926, 0.09445201,
0.13371421, 0.16924426, 0.16342211, 0.13938707, 0.12998206,
0.15834639, 0.14789637, 0.10296131, 0.11774062, 0.13535635,
0.10549917, 0.11296347, 0.14028279, 0.10340917, 0.09967702,
0.12281634, 0.16431782, 0.15192709, 0.15088209, 0.11744204,
0.11072418, 0.20492359, 0.09101844, 0.09639273, 0.09654201,
0.12445848, 0.1592421 , 0.14909066, 0.09893059, 0.11818847,
0.13610278, 0.14998637, 0.12520491, 0.15879424, 0.16207853,
0.11191846, 0.15700281, 0.17073712, 0.10281202, 0.09280987,
0.17640998, 0.10102059, 0.16595997, 0.16416854, 0.2122386 ,
0.11221704, 0.19462286, 0.14401494, 0.15849567, 0.18402356,
```

Question-14:

Measure the performance using Metrics.

Solution:

```
In [63]: from sklearn.metrics import r2_score  
acc=r2_score(y_test,y_pred)  
acc
```

Out[63]: 0.7680363609738088

```
In [64]: from sklearn import metrics  
np.sqrt(metrics.mean_squared_error(y_test,y_pred))
```

Out[64]: 0.01811026443738919

```
In [65]: regressor.predict([[1034]])
```

Out[65]: array([33.74303656])