ASSIGNMENT -2
DATA VISUALIZATION AND DATA PRE-PROCESSING

Assignment Date	17 September 2022
Student Name	Mohana Sowdesh R
Student Roll Number	727719EUCS091
Maximum Marks	2 Marks

Question-1:

Download the dataset: Dataset

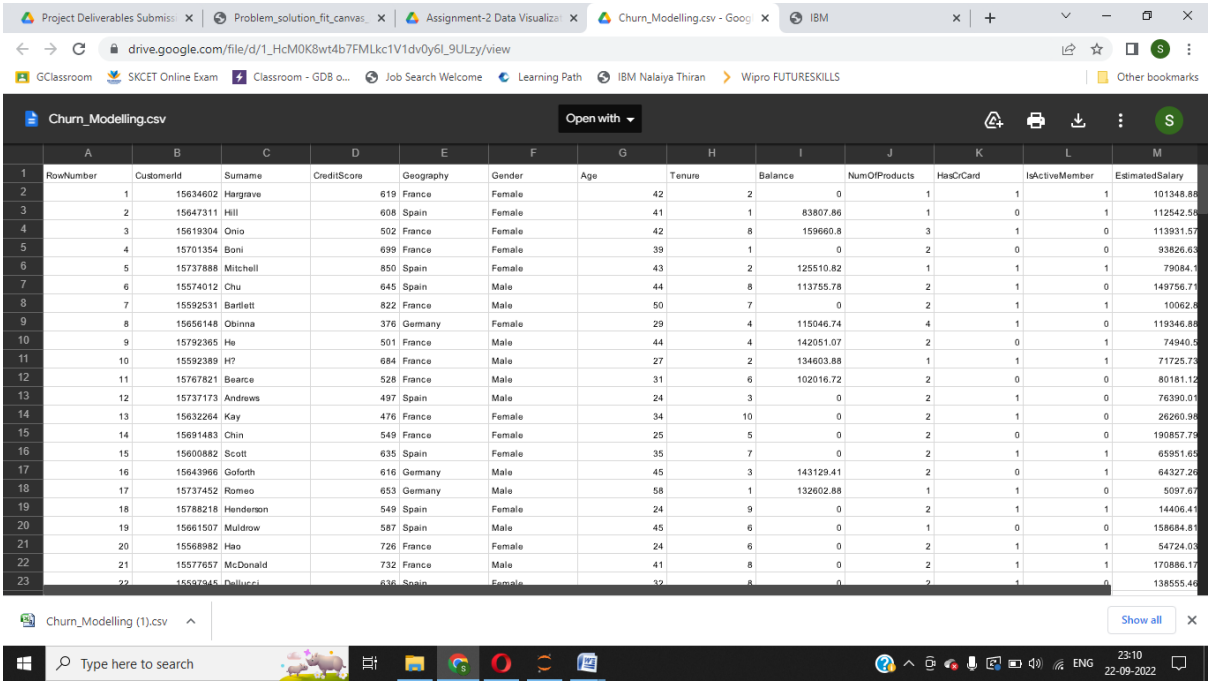

Question-2:

Load the dataset.

Solution:

```
import pandas as pd

import numpy as np

import matplotlib.pyplot as plt

import seaborn as sns

df = pd.read_csv("C://Users\Mohana Sowdesh//Desktop//IBM Nalaiya Thiran// Dataset//
Churn_ Modelling.csv")
df.head()
```

```
In [2]: import pandas as pd
        import numpy as np
        import matplotlib.pyplot as plt
        import seaborn as sns
```

```
In [4]: df = pd.read_csv("C://Users\Mohana Sowdesh//Desktop//IBM Nalaiya Thiran//Dataset//Churn_Modelling.csv")
        df.head()
```

Out[4]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.88
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	113931.57
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	93826.63
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10

Question-3:

Perform Below Visualizations.

● Univariate Analysis

```
In [5]: sns.distplot(df.EstimatedSalary)
```

```
C:\Users\Mohana Sowdesh\anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated fu
nction and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with si
milar flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```

```
Out[5]: <AxesSubplot:xlabel='EstimatedSalary', ylabel='Density'>
```

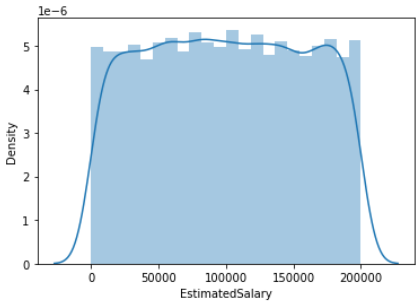

● Bi - Variate Analysis

```
In [8]: sns.scatterplot(df.CreditScore,df.EstimatedSalary)
```

```
Out[8]: <AxesSubplot:xlabel='CreditScore', ylabel='EstimatedSalary'>
```

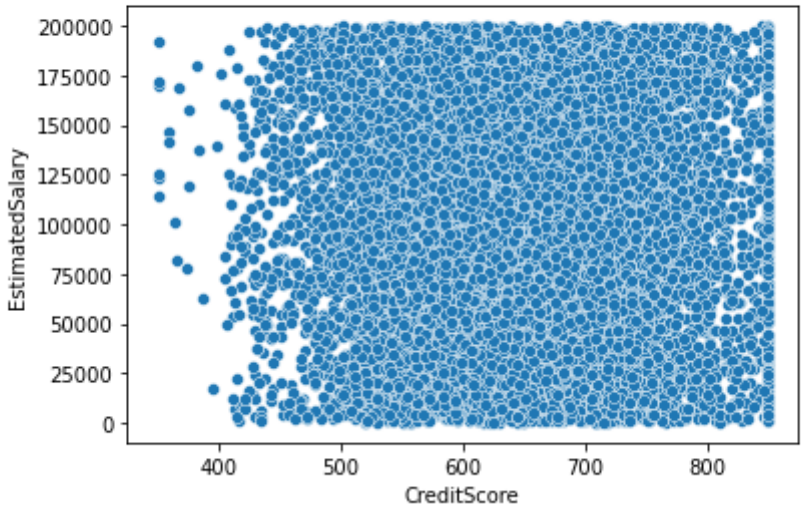

● Multi - Variate Analysis

```
In [9]: sns.pairplot(df)
Out[9]: <seaborn.axisgrid.PairGrid at 0x13f5bf2ab20>
```

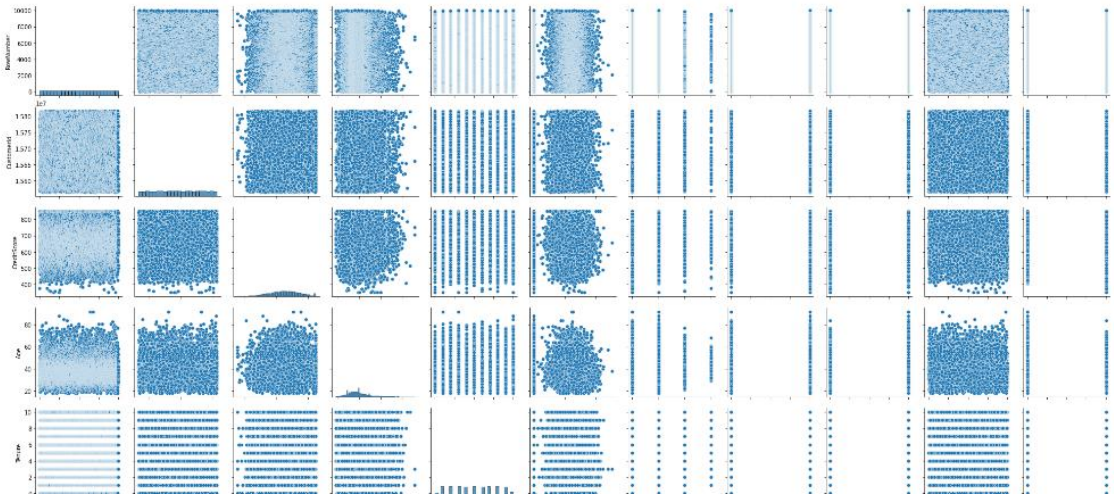

Question-4:

Perform descriptive statistics on the dataset.

```
In [9]: df['Age'].mean()
Out[9]: 38.9218

In [11]: df['NumOfProducts'].median()
Out[11]: 1.0

In [12]: df['HasCrCard'].mode()
Out[12]: 0    1
         dtype: int64

In [13]: df.skew()
Out[13]: RowNumber          0.000000
         CustomerId         0.001149
         CreditScore       -0.071607
         Age                1.011320
         Tenure             0.010991
         Balance           -0.141109
         NumOfProducts      0.745568
         HasCrCard         -0.901812
         IsActiveMember    -0.060437
         EstimatedSalary    0.002085
         Exited             1.471611
         dtype: float64
```

In [14]:
```
df.kurt()
```

Out[14]:
```
RowNumber         -1.200000
CustomerId        -1.196113
CreditScore       -0.425726
Age                1.395347
Tenure            -1.165225
Balance           -1.489412
NumOfProducts      0.582981
HasCrCard         -1.186973
IsActiveMember    -1.996747
EstimatedSalary   -1.181518
Exited             0.165671
dtype: float64
```

In [19]:
```
df.var()
```

Out[19]:
```
RowNumber          8.334167e+06
CustomerId         5.174815e+09
CreditScore        9.341860e+03
Age                1.099941e+02
Tenure             8.364673e+00
Balance            3.893436e+09
NumOfProducts      3.383218e-01
HasCrCard          2.077905e-01
IsActiveMember     2.497970e-01
EstimatedSalary    3.307457e+09
Exited             1.622225e-01
dtype: float64
```

In [20]:
```
df.std()
```

Out[20]:
```
RowNumber          2886.895680
CustomerId        71936.186123
CreditScore          96.653299
Age                  10.487806
Tenure                2.892174
Balance           62397.405202
NumOfProducts         0.581654
HasCrCard             0.455840
IsActiveMember        0.499797
EstimatedSalary   57510.492818
Exited                0.402769
dtype: float64
```

Question-5:

Handle the Missing values.

```
In [21]: df.isna().any()

Out[21]: RowNumber          False
         CustomerId         False
         Surname            False
         CreditScore        False
         Geography          False
         Gender             False
         Age                False
         Tenure             False
         Balance            False
         NumOfProducts      False
         HasCrCard          False
         IsActiveMember     False
         EstimatedSalary    False
         Exited             False
         dtype: bool
```

```
In [22]: df.isna().sum()

Out[22]: RowNumber          0
         CustomerId         0
         Surname            0
         CreditScore        0
         Geography          0
         Gender             0
         Age                0
         Tenure             0
         Balance            0
         NumOfProducts      0
         HasCrCard          0
         IsActiveMember     0
         EstimatedSalary    0
         Exited             0
         dtype: int64
```

```
In [24]: df['Age'].fillna(df['Age'].mean(),inplace=True)
         df
```

Out[24]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSa
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	10134
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	11254
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	11393
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	9382
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	7908
...	...	...	...	...	...	...	...	...	...	...	...	...	...
9995	9996	15606229	Obijiaku	771	France	Male	39	5	0.00	2	1	0	9627
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61	1	1	1	10169
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1	0	1	4208
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	75075.31	2	1	0	9288
9999	10000	15628319	Walker	792	France	Female	28	4	130142.79	1	1	0	3819

10000 rows × 14 columns

Question-6:

Find the outliers and replace the outliers

```
In [43]: Q1=df.Age.quantile(0.25)
         Q2=df.Age.quantile(0.75)
         IQR=Q2-Q1
         print(IQR)
```

```
12.0
```

```
In [44]: df=df[~((df.Age<(Q1-1.5*IQR))|(df.Age>(Q2+1.5*IQR)))]
         df
```

Out[44]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSa
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	10134
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	11254
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	11393
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	9382
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	7908
...	...	...	...	...	...	...	...	...	...	...	...	...	...
9995	9996	15606229	Obijiaku	771	France	Male	39	5	0.00	2	1	0	9627
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61	1	1	1	10169
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1	0	1	4208
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	75075.31	2	1	0	9288
9999	10000	15628319	Walker	792	France	Female	28	4	130142.79	1	1	0	3819

9641 rows × 14 columns

Question-7:

Check for Categorical columns and perform encoding.

```
In [46]: df['HasCrCard'].replace({1:'Yes',0:'No'},inplace=True)
         df.head()
```

Out[46]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	Yes	1	101348.88
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	No	1	112542.58
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	Yes	0	113931.57
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	No	0	93826.63
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	Yes	1	79084.10

Question-8:

Split the data into dependent and independent variables.

```
In [50]: data_main= pd.get_dummies(df,columns=['Gender'])
         data_main
```

Out[50]:

Surname	CreditScore	Geography	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited	Gender_Female	Gender_Male
Hargrave	619	France	42	2	0.00	1	Yes	1	101348.88	1	1	0
Hill	608	Spain	41	1	83807.86	1	No	1	112542.58	0	1	0
Onio	502	France	42	8	159660.80	3	Yes	0	113931.57	1	1	0
Boni	699	France	39	1	0.00	2	No	0	93826.63	0	1	0
Mitchell	850	Spain	43	2	125510.82	1	Yes	1	79084.10	0	1	0
...	...	...	...	...	...	...	...	...	...	...	...	...
Obijiaku	771	France	39	5	0.00	2	Yes	0	96270.64	0	0	1
Johnstone	516	France	35	10	57369.61	1	Yes	1	101699.77	0	0	1
Liu	709	France	36	7	0.00	1	No	1	42085.58	1	1	0
Sabbatini	772	Germany	42	3	75075.31	2	Yes	0	92888.52	1	0	1
Walker	792	France	28	4	130142.79	1	Yes	0	38190.78	0	1	0

```
In [51]: y = data_main['Tenure']
         y

Out[51]: 0        2
         1        1
         2        8
         3        1
         4        2
                 ..
         9995     5
         9996    10
         9997     7
         9998     3
         9999     4
         Name: Tenure, Length: 9641, dtype: int64
```

```
In [52]: x = data_main.drop(columns='Tenure',axis=1)
         x.head()
```

Out[52]:

tomerId	Surname	CreditScore	Geography	Age	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited	Gender_Female	Gender_Male
i634602	Hargrave	619	France	42	0.00	1	Yes	1	101348.88	1	1	0
i647311	Hill	608	Spain	41	83807.86	1	No	1	112542.58	0	1	0
i619304	Onio	502	France	42	159660.80	3	Yes	0	113931.57	1	1	0
i701354	Boni	699	France	39	0.00	2	No	0	93826.63	0	1	0
i737888	Mitchell	850	Spain	43	125510.82	1	Yes	1	79084.10	0	1	0

Question-9:

Scale the independent variables

```
In [55]: x=df.iloc[:,6:7].values
         from sklearn.preprocessing import StandardScaler
         std=StandardScaler()
         x=std.fit_transform(x)
         x

Out[55]: array([[ 0.47806838],
                [ 0.36446646],
                [ 0.47806838],
                ...,
                [-0.20354316],
                [ 0.47806838],
                [-1.11235856]])
```

Question-10:

Split the data into training and testing

```
In [56]: from sklearn.model_selection import train_test_split
         x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=0)
```

```
In [57]: x_train
```

```
Out[57]: array([[-2.24837781],
                [ 0.59167031],
                [ 1.04607801],
                ...,
                [-0.54434894],
                [ 1.04607801],
                [-0.43074701]])
```

```
In [58]: x_test
```

```
Out[58]: array([[ 1.50048571],
                [-0.20354316],
                [ 0.36446646],
                ...,
                [ 0.81887416],
                [-0.88515471],
                [ 0.13726261]])
```

```
In [59]: y_train
```

```
Out[59]: 746     2
         1788    8
         1057    1
         7559    2
         1141    5
                ..
         8184    3
         9567    4
         5042    3
         3370    6
         2819    5
         Name: Tenure, Length: 7712, dtype: int64
```

```
In [60]: y_test
```

```
Out[60]: 2454    1
         944     8
         3938    1
         4109    1
         8573    8
                ..
         1465    3
         8409    9
         5624    1
         2817    8
         6851    9
         Name: Tenure, Length: 1929, dtype: int64
```